



A dark blue background featuring a complex network of interconnected nodes (purple, teal, and grey dots) and lines, resembling a neural network or a circuit board. Interspersed among the nodes are binary digits (0s and 1s) in various colors, suggesting data flow or processing. The overall theme is technology, data, and connectivity.

Build. Unify. Scale.

WIFI SSID:SparkAISummit | Password: UnifiedAnalytics

ORGANIZED BY
 **databricks**



Spark Core – Proper Optimization

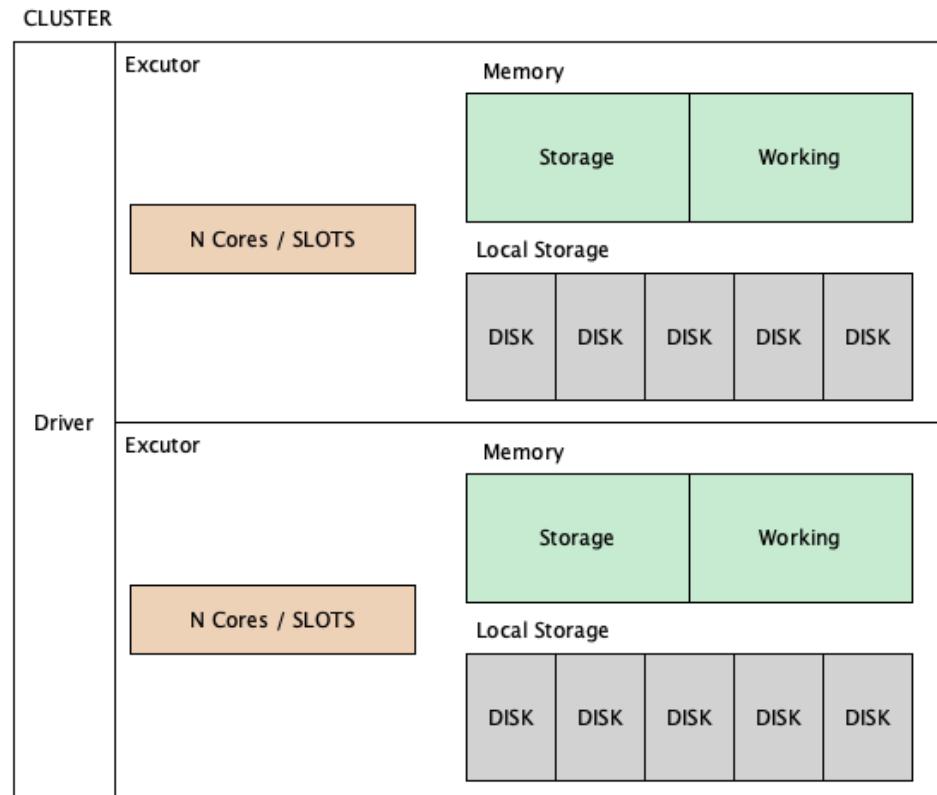
Daniel Tomes, Databricks

#UnifiedAnalytics #SparkAIsummit

Talking Points

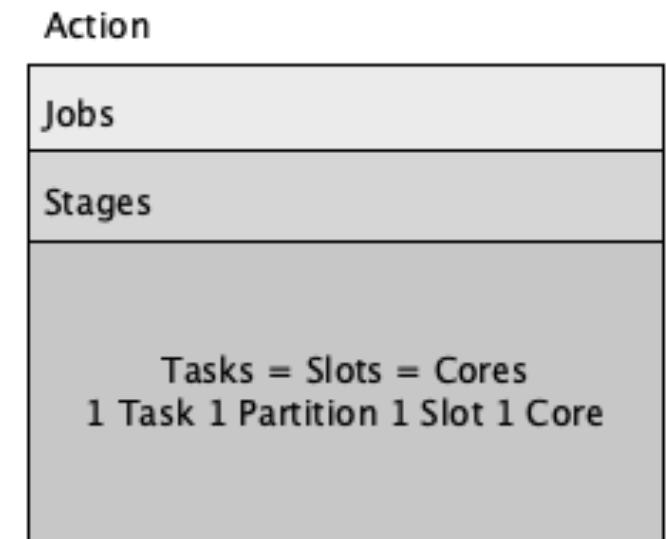
- Spark Hierarchy
- The Spark UI
- Rightsizing & Optimizing
- Advanced Optimizations

Spark Hierarchy



Spark Hierarchy

- Actions are eager
 - Made of transformations (lazy)
 - narrow
 - wide (requires shuffle)
 - Spawn jobs
 - Spawn Stages
 - Spawn Tasks
 - » Do work & utilize hardware



Navigating The Spark UI

DEMO

Understand Your Hardware

- Core Count & Speed
- Memory Per Core (Working & Storage)
- Local Disk Type, Count, Size, & Speed
- Network Speed & Topology
- Data Lake Properties (rate limits)
- Cost / Core / Hour
 - Financial For Cloud
 - Opportunity for Shared & On Prem

Get A Baseline

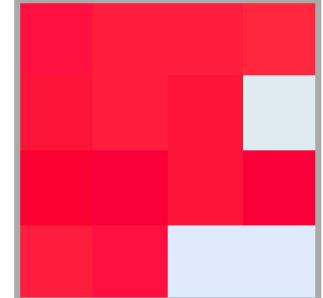
- Is your action efficient?
 - Long Stages, Spills, Laggard Tasks, etc?
- CPU Utilization
 - GANGLIA / YARN / Etc
 - Tails

Goal

CPUs Total: **216**
Hosts up: **14**
Hosts down: **0**

Current Load Avg (15, 5, 1m):
34%, 79%, 125%
Avg Utilization (last hour):
0%

Server Load Distribution



Minimize Data Scans (Lazy Load)

- Data Skipping
 - HIVE Partitions
 - Bucketing
 - Only Experts – Nearly Impossible to Maintain
 - Databricks Delta Z-Ordering
 - [What is It](#)
 - [How To Do It](#)

Cmd 8

```
1 val master_no_lazy = ss_sub
2   .join(dt_sub.withColumnRenamed("d_date_sk", "ss_sold_date_sk"), Seq("ss_sold_date_sk"))
3   .join(item.withColumnRenamed("i_item_sk", "ss_item_sk"), Seq("ss_item_sk"))
4   .join(inv.withColumnRenamed("inv_item_sk", "ss_item_sk"), Seq("ss_item_sk"))
```

No Lazy Loading

Cmd 8

```

1 val master_no_lazy = ss_sub
2 .join(dt_sub.withColumnRenamed("d_date_sk", "ss_sold_date_sk"), Seq("ss_sold_date_sk"))
3 .join(item.withColumnRenamed("i_item_sk", "ss_item_sk"), Seq("ss_item_sk"))
4 .join(inv.withColumnRenamed("inv_item_sk", "ss_item_sk"), Seq("ss_item_sk"))

```

Simple

Active Stages (1)

Stage Id	Pool Name	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
6	7127041251392214025	master_no_lazy.write.mode("overwrite").format("... save at command-885910:1 +details (kill)	2019/03/31 17:11:16	6.2 h	0/200 (96 running)		58.7 GB		

Extra Shuffle Partitions

Active Stages (1)

Stage Id	Pool Name	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
6	3026732521940153177	spark.conf.set("spark.sql.shuffle.partitions", ... save at command-885910:2 +details (kill)	2019/03/31 23:30:54	10 min	0/2016 (96 running)		5.8 GB		

With Lazy Loading

Cmd 12

```

1 val master_lazy = ss_sub
2 .join(dt_sub
3   .filter(year('d_date).between(2000,2001)),
4   'd_date_sk === 'ss_sold_date_sk')
5 .join(item.withColumnRenamed("i_item_sk", "ss_item_sk"), Seq("ss_item_sk"))
6 .join(inv,
7   'inv_item_sk === 'ss_item_sk &&
8   'inv_date_sk === 'ss_sold_date_sk)

```

Details for Stage 21 (Attempt 0)

Total Time Across All Tasks: 12.9 h
 Locality Level Summary: Process local: 200
 Output: 21.8 GB / 3112776340
 Shuffle Read: 53.9 GB / 1883116073
 Shuffle Spill (Memory): 552.0 GB
 Shuffle Spill (Disk): 67.4 GB

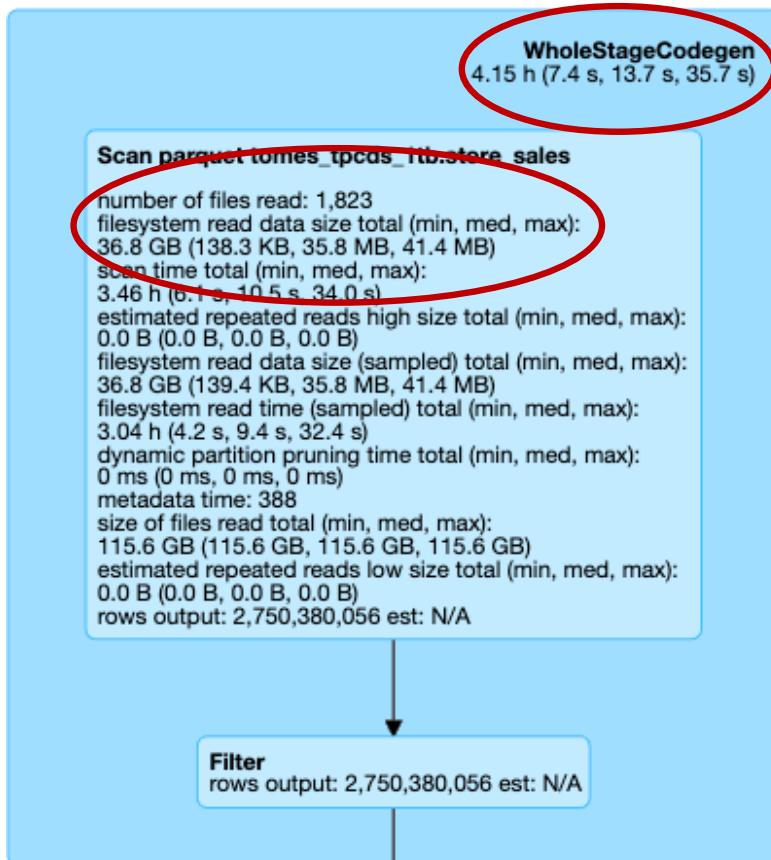
- ▶ DAG Visualization
- ▶ Show Additional Metrics
- ▶ Event Timeline

Summary Metrics for 200 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	1.7 min	3.8 min	3.9 min	4.2 min	4.6 min
GC Time	2 s	7 s	11 s	23 s	30 s
Output Size / Records	111.0 MB / 15333760	111.6 MB / 15524100	111.7 MB / 15566840	111.9 MB / 15608700	114.0 MB / 15771320
Shuffle Read Size / Records	275.0 MB / 9379740	275.8 MB / 9405412	276.1 MB / 9415012	276.4 MB / 9425519	277.7 MB / 9458918
Shuffle spill (memory)	0.0 B	2.9 GB	2.9 GB	2.9 GB	2.9 GB
Shuffle spill (disk)	0.0 B	358.9 MB	359.7 MB	360.1 MB	360.6 MB

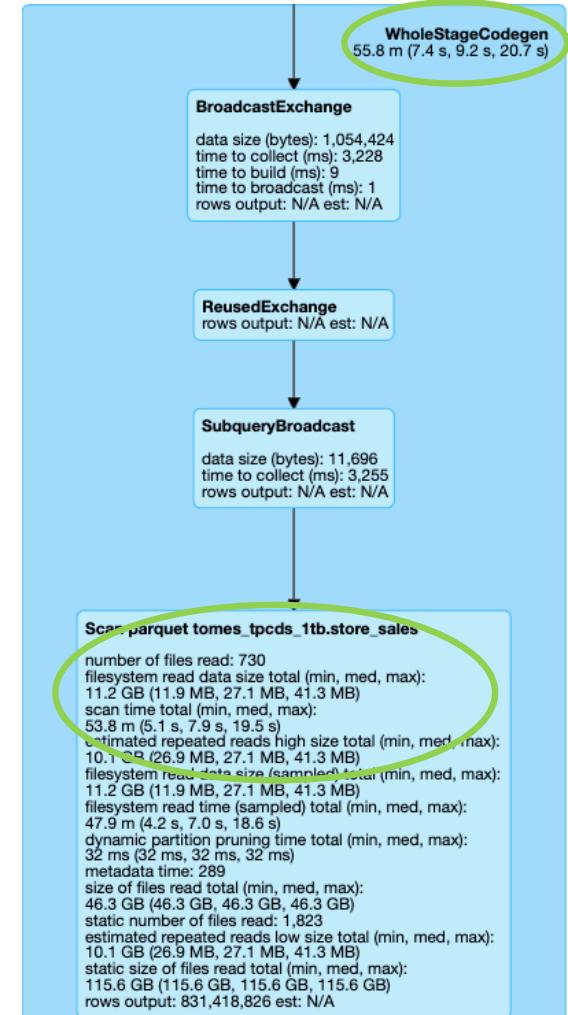
Stage Id	Pool Name	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
21	238012049243624904	master_lazy_item_bc.write.format("parquet").sav... save at command-885895:1 +details	2019/03/31 14:17:44	8.5 min	200/200		21.8 GB	53.9 GB	
20	238012049243624904	master_lazy_item_bc.write.format("parquet").sav... save at command-885895:1 +details	2019/03/31 14:16:15	27 s	121/121		3.8 GB		8.6 GB
19	238012049243624904	master_lazy_item_bc.write.format("parquet").sav... save at command-885895:1 +details	2019/03/31 14:16:15	1.3 min	452/452		14.7 GB		45.4 GB

Without Partition Filter



Shrink Partition Range Using a Filter on HIVE Partitioned Column

With Partition Filter



Partitions – Definition

Each of a number of portions into which some operating systems divide memory or storage

~~HIVE PARTITION == SPARK PARTITION~~

Spark Partitions – Types

- **Input**
 - Controls - Size
 - spark.default.parallelism (don't use)
 - spark.sql.files.maxPartitionBytes (mutable)
 - assuming source has sufficient partitions
- **Shuffle**
 - Control = Count
 - spark.sql.shuffle.partitions
- **Output**
 - Control = Size
 - Coalesce(n) to shrink
 - Repartition(n) to increase and/or balance (shuffle)
 - df.write.option("maxRecordsPerFile", N)

Partitions – Shuffle – Default

Default = 200 Shuffle Partitions

Partitions – Right Sizing – Shuffle – Master Equation

- Largest Shuffle Stage
 - Target Size \leq 200 MB/partition
- Partition Count = Stage Input Data / Target Size
 - Solve for Partition Count

EXAMPLE

Shuffle Stage Input = 210GB

$$x = 210000\text{MB} / 200\text{MB} = 1050$$

`spark.conf.set("spark.sql.shuffle.partitions", 1050)`

BUT -> If cluster has 2000 cores

`spark.conf.set("spark.sql.shuffle.partitions", 2000)`

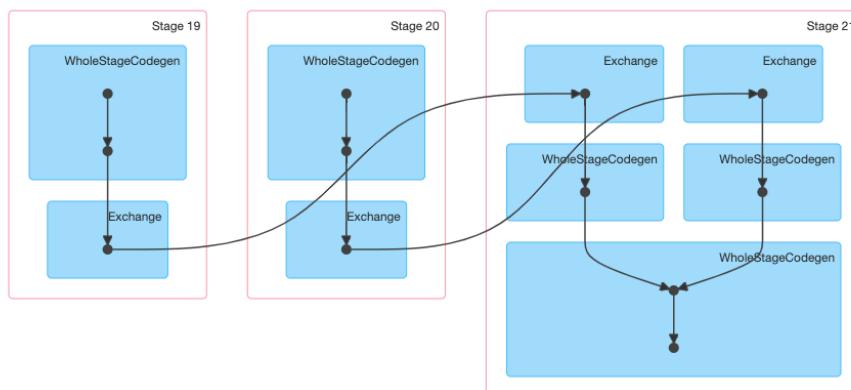
Cluster Spec

96 cores @ 7.625g/core
 3.8125g Working Mem
 3.8125g Storage Mem

Details for Job 11

Status: SUCCEEDED
 Associated SQL Query: 95
 Job Group: 238012049243624904_7941856623374133785_21481c00c02f4ff2aeee8f2d8342a3ea5
 Completed Stages: 3

- ▶ Event Timeline
- ▶ DAG Visualization



Completed Stages (3)

Stage Id	Pool Name	Description	Submitted	Duration	Tasks: Succeeded / Total	Input	Output	Shuffle Read	Shuffle Write
21	238012049243624904	master_lazy_item_bc.write.format("parquet").sav... save at command-885895:1	2019/03/31 14:17:34 +details	8.5 min	200/200	21.8 GB	53.9 GB		
20	238012049243624904	master_lazy_item_bc.write.format("parquet").sav... save at command-885895:1	2019/03/31 14:16:15 +details	27 s	131/131	3.8 GB			8.6 GB
19	238012049243624904	master_lazy_item_bc.write.format("parquet").sav... save at command-885895:1	2019/03/31 14:16:15 +details	1.3 min	452/452	14.7 GB			45.4 GB

Stage 21 -> Shuffle Fed By Stage 19 & 20

THUS

Stage 21 Shuffle Input = 45.4g + 8.6g == 54g

Default Shuffle Partition == 200 == 54000mb/200parts ==~ 270mb/shuffle part

Details for Stage 21 (Attempt 0)

Total Time Across All Tasks: 12.9 h
 Locality Level Summary: Process local: 200
 Output: 21.8 GB / 3112776340
 Shuffle Read: 53.9 GB / 1883116073
 Shuffle Spill (Memory): 552.0 GB
 Shuffle Spill (Disk): 67.4 GB

▶ DAG Visualization
 ▶ Show Additional Metrics
 ▶ Event Timeline
 Summary Metrics for 200 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	1.7 min	3.8 min	3.9 min	4.2 min	4.6 min
GC Time	2 s	7 s	11 s	23 s	30 s
Output Size / Records	111.0 MB / 15333760	111.6 MB / 15324100	111.7 MB / 15566840	111.9 MB / 15608700	114.0 MB / 15771320
Shuffle Read Size / Records	275.0 MB / 9379740	275.8 MB / 9405412	276.1 MB / 9415012	276.4 MB / 9425519	277.7 MB / 9458918
Shuffle spill (memory)	0.0 B	2.9 GB	2.9 GB	2.9 GB	2.9 GB
Shuffle spill (disk)	0.0 B	358.9 MB	359.7 MB	360.1 MB	360.6 MB

Spills

Cluster Spec

96 cores @ 7.625g/core
3.8125g Working Mem
3.8125g Storage Mem

```
1 spark.conf.set("spark.sql.shuffle.partitions", 480)
```

Completed Stages (3)

Stage Id	Pool Name	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
37	238012049243624904	spark.conf.set("spark.sql.shuffle.partitions", ... save at command-885895:2 +details)	2019/03/31 14:53:45	7.8 min	480/480		19.9 GB	54.0 GB	
36	238012049243624904	spark.conf.set("spark.sql.shuffle.partitions", ... save at command-885895:2 +details)	2019/03/31 14:52:25	25 s	131/131	3.8 GB			8.7 GB
35	238012049243624904	spark.conf.set("spark.sql.shuffle.partitions", ... save at command-885895:2 +details)	2019/03/31 14:52:25	1.3 min	452/452	14.7 GB			45.3 GB

480 shuffle partitions – WHY?

Target shuffle part size == 100m
 $p = 54g / 100m = 540$

$540p / 96 \text{ cores} = 5.625$
 $96 * 5 = 480$

If $p == 540$ another 60p have to be loaded and processed after first cycle is complete

Details for Stage 37 (Attempt 0)

Total Time Across All Tasks: 11.8 h
Locality Level Summary: Process local: 480
Output: 19.9 GB / 3112776340
Shuffle Read: 54.0 GB / 1883116073

- ▶ DAG Visualization
- ▶ Show Additional Metrics
- ▶ Event Timeline

NO SPILL

Summary Metrics for 480 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	1.2 min	1.5 min	1.5 min	1.6 min	1.8 min
GC Time	0.8 s	2 s	2 s	3 s	12 s
Output Size / Records	41.5 MB / 6362420	42.2 MB / 6456400	42.5 MB / 6484840	42.7 MB / 6515940	43.4 MB / 6601140
Shuffle Read Size / Records	114.3 MB / 3895329	115.0 MB / 3917334	115.2 MB / 3923186	115.4 MB / 3929412	115.9 MB / 3948584

Input Partitions – Right Sizing

- Use Spark Defaults (128MB) unless...
 - Increase Parallelism
 - Heavily Nested/Repetitive Data
 - Generating Data – i.e. Explode
 - Source Structure is not optimal (upstream)
 - UDFs

```
spark.conf.set("spark.sql.files.maxPartitionBytes", 16777216)
```

Cmd 48

```
1 spark.conf.set("spark.sql.files.maxPartitionBytes", 134217728) 128mb
2 val master_source_10p = spark.read.parquet("/tmp/tomes/scratch/SAIS19/10p_source")
3 master_source_10p.rdd.partitions.size
```

▶ (1) Spark Jobs

▶ [master_source_10p: org.apache.spark.sql.DataFrame = [ss_item_sk: integer, ss_customer_sk: integer ... 33 m
master_source_10p: org.apache.spark.sql.DataFrame = [ss_item_sk: int, ss_customer_sk: in
res34: Int = 90]

```
1 spark.conf.set("spark.sql.files.maxPartitionBytes", 1024 * 1024 * 16) 16mb
2 val master_source_10p = spark.read.parquet("/tmp/tomes/scratch/SAIS19/10p_source")
3 master_source_10p.rdd.partitions.size
```

▶ (1) Spark Jobs

▶ [master_source_10p: org.apache.spark.sql.DataFrame = [ss_item_sk: integer, ss_customer_sk: integer ... 33 r
master_source_10p: org.apache.spark.sql.DataFrame = [ss_item_sk: int, ss_customer_sk: i
res36: Int = 710]

Stage Id ▾	Pool Name	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output
3	7427136891458074341	spark.conf.set("spark.sql.files.maxPartitionBy... save at command-899638:4 +details	2019/04/23 13:49:35	4.6 min	90/90	11.9 GB	24.1 GB
<hr/>							
Metric	Min	25th percentile	Median	75th percentile	Max		
Duration	2.9 min	3.8 min	4.4 min	4.5 min	4.5 min		
GC Time	1 s	1 s	1 s	2 s	2 s		
Input Size / Records	105.9 MB / 26344585	138.9 MB / 35511851	139.0 MB / 35544755	139.1 MB / 35576807	139.4 MB / 35647533		
Output Size / Records	213.8 MB / 26344585	280.5 MB / 35511851	281.1 MB / 35544755	281.5 MB / 35576807	282.6 MB / 35647533		

Metric	Min	25th percentile	Median	75th percentile	Max		
Duration	6 s	1.1 min	1.5 min	1.6 min	2.5 min		
GC Time	48 ms	3 s	4 s	5 s	6 s		
Input Size / Records	11.0 MB / 0	26.9 MB / 4414822	27.0 MB / 4441791	27.1 MB / 4471397	27.4 MB / 4537083		
Output Size / Records	0.0 B / 0	25.1 MB / 4414822	25.4 MB / 4441791	25.6 MB / 4471397	26.2 MB / 4537083		

Stage Id ▾	Pool Name	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output
1	7427136891458074341	spark.conf.set("spark.sql.files.maxPartitionBy... save at command-899639:4 +details	2019/04/23 13:46:34	2.9 min	710/710	18.6 GB	17.4 GB

Output Partitions – Right Sizing

- Write Once -> Read Many
 - More Time to Write but Faster to Read
- Perfect writes limit parallelism
 - Compactions (minor & major)

Write Data Size = 14.7GB

Desired File Size = 1500MB

Max write stage parallelism = 10

96 – 10 == 86 cores idle during write

Only 10 Cores Used

Completed Stages (1)

Stage Id	Pool Name	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
113	9020956004722677074	val item_bc = broadcast(item.withColumnRenamed(... save at command-885938:10 +details	2019/03/31 19:51:10	14 min	10/10	18.9 GB	14.3 GB		

Cmd 21

```
1 getSizeInfo(s"$s3Prefix/store_sales_2y_10p")
```

▶ (4) Spark Jobs

Table store_sales_2y_10p has 10 files and average 1.5304093662 gb with stddev of 250 mb

Command took 0.61 seconds -- by daniel.tomes@databricks.com at 3/31/2019, 2:21:49 PM on Tomes_Base

Average File Size == 1.5g

All 96 Cores Used

Completed Stages (1)

Stage Id	Pool Name	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
97	9020956004722677074	val item_bc = broadcast(item.withColumnRenamed(... save at command-885971:10 +details	2019/03/31 19:43:41	1.8 min	96/96	15.1 GB	14.3 GB		

Cmd 22

```
1 getSizeInfo(s"$s3Prefix/store_sales_2y_96p")
```

▶ (4) Spark Jobs

Table store_sales_2y_10p has 96 files and average 0.15996809254166666 gb with stddev of 35 mb

Command took 2.00 seconds -- by daniel.tomes@databricks.com at 3/31/2019, 3:46:45 PM on Tomes_Base

Average File Size == 0.16g

Output Partitions – Composition

- `df.write.option("maxRecordsPerFile", n)`
- `df.coalesce(n).write...`
- `df.repartition(n).write...`
- `df.repartition(n, [colA, ...]).write...`
- `spark.sql.shuffle.partitions(n)`
- `df.localCheckpoint(...).repartition(n).write...`
- `df.localCheckpoint(...).coalesce(n).write...`

Partitions – Why So Serious?

- Avoid The Spill
- Maximize Parallelism
 - Utilize All Cores
 - Provision only the cores you need

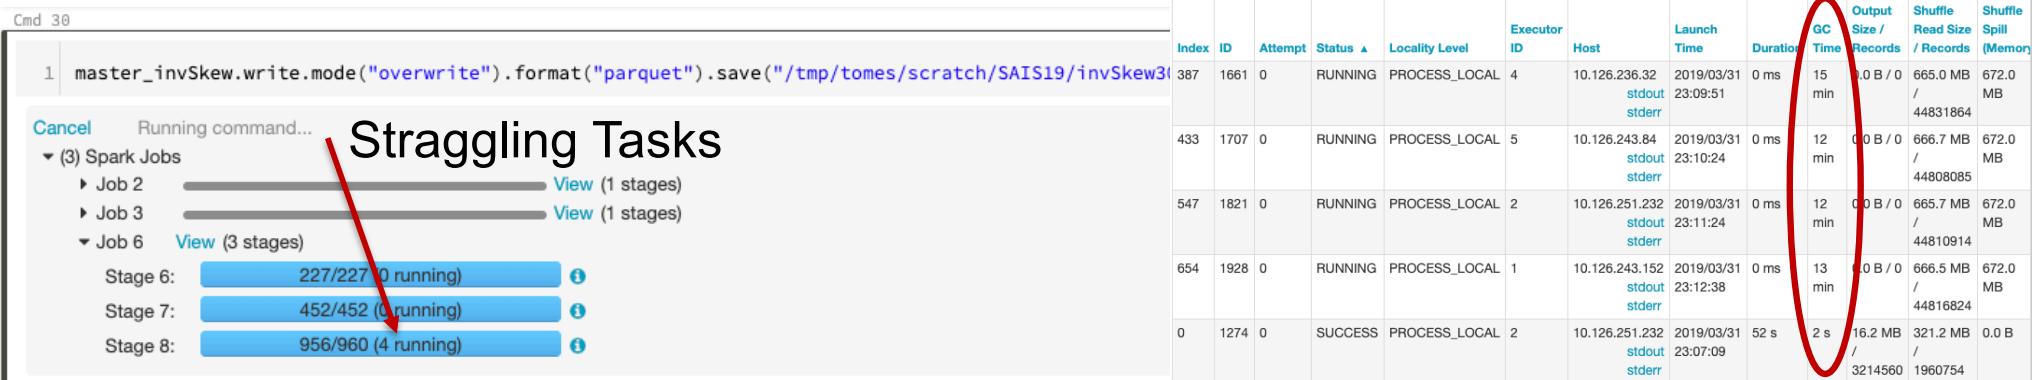
Advanced Optimizations

- Finding Imbalances
- Persisting
- Join Optimizations
- Handling Skew
- Expensive Operations
- UDFs
- Multi-Dimensional Parallelism

Balance

- Maximizing Resources Requires Balance
 - Task Duration
 - Partition Size
- SKEW
 - When some partitions are significantly larger than most

Input Partitions
Shuffle Partitions
Output Files
Spills
GC Times



75th percentile ~ 2m recs

max ~ 45m recs

stragglers take > 22X longer IF no spillage

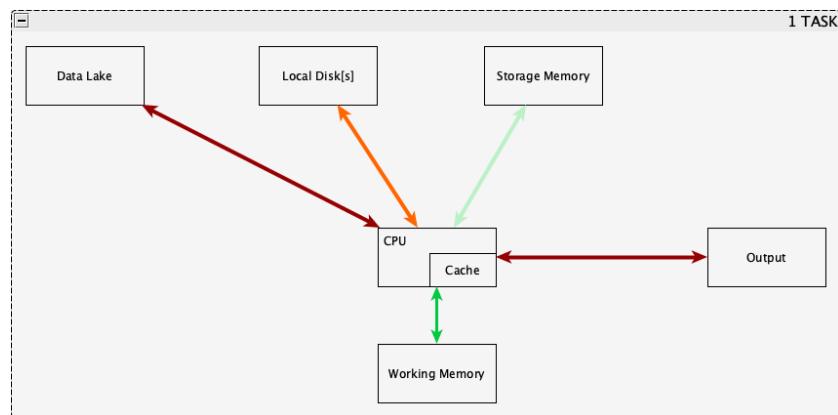
With spillage, 100Xs longer

Summary Metrics for 956 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	0 ms	40 s	47 s	52 s	3.7 min
GC Time	0.3 s	0.8 s	1 s	4 s	3.2 min
Output Size / Records	0.0 B / 0	16.1 MB / 3219340	16.3 MB / 3242100	16.5 MB / 3263540	17.0 MB / 3323340
Shuffle Read Size / Records	317.2 MB / 1940846	320.7 MB / 1957433	321.6 MB / 1961754	322.4 MB / 1965647	666.7 MB / 44831864
Shuffle spill (memory)	0.0 B	0.0 B	0.0 B	0.0 B	672.0 MB
Shuffle spill (disk)	0.0 B	0.0 B	0.0 B	0.0 B	91.5 MB

Minimize Data Scans (Persistence)

- Persistence
 - Not Free
- Repetition
 - SQL Plan

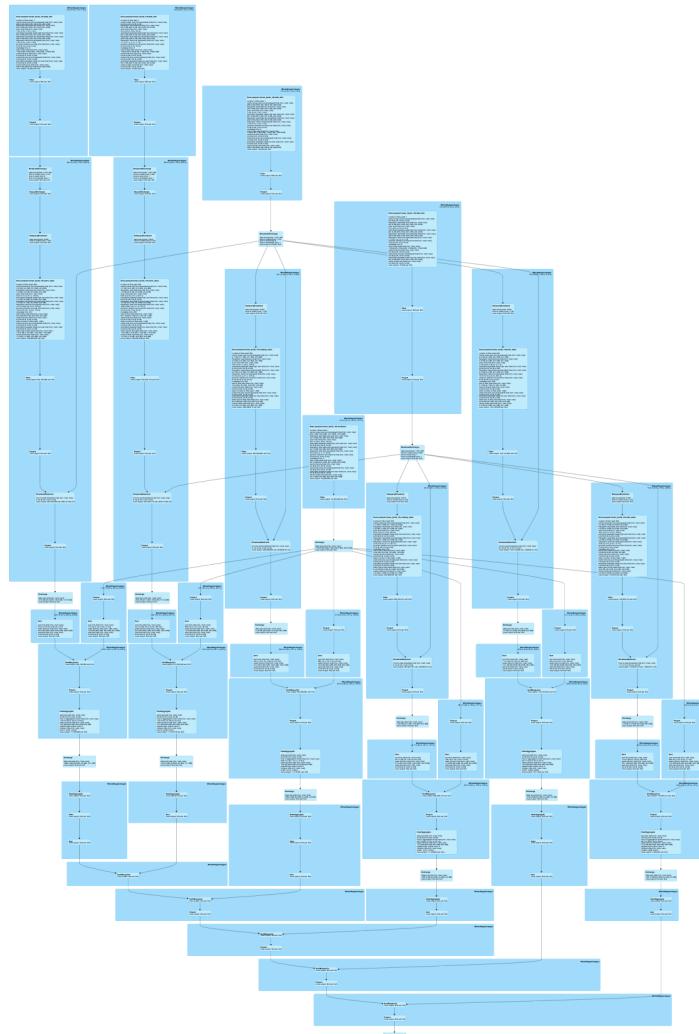


`df.cache == df.persist(StorageLevel.MEMORY_AND_DISK)`

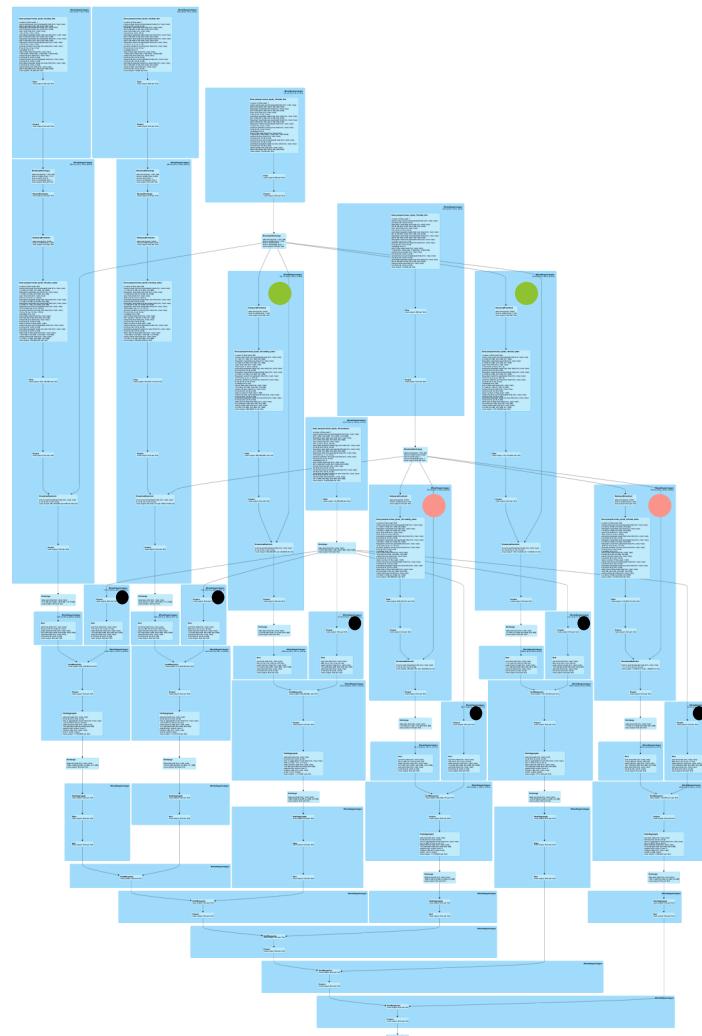
- Types

- Default (MEMORY_AND_DISK)
 - Deserialized
- Deserialized = Faster = Bigger
- Serialized = Slower = Smaller
- _2 = Safety = 2X bigger
- MEMORY_ONLY
- DISK_ONLY

Don't Forget To Cleanup!
`df.unpersist`



TPCDS Query 4



Minimize Data Scans (Delta Cache)

CACHE SELECT column_name[, column_name, ...] FROM [db_name.]table_name [WHERE boolean_expression]

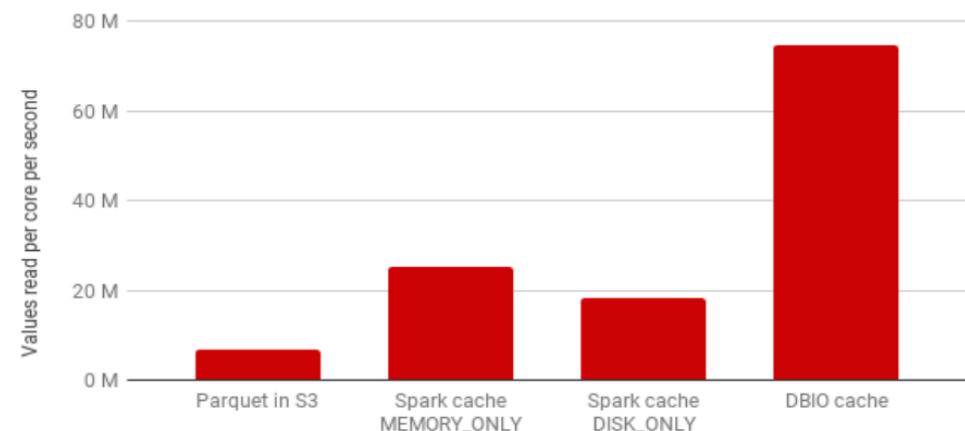
Jobs	Stages	Storage	Environment	Executors	SQL	JDBC/ODBC Server
Storage						
Parquet IO Cache						
Host	Disk Usage	Max Disk Usage Limit	Percent Disk Usage	Metadata Cache Size	Max Metadata Cache Size Limit	Percent Metadata Usage
10.231.249.153	168.1 GB	1855.5 GB	9 %	5.8 MB	37.1 GB	0 %
10.231.250.190	175.4 GB	1855.5 GB	9 %	6.0 MB	37.1 GB	0 %
Total	343.5 GB	3.6 TB	9 %	11.8 MB	74.2 GB	0 %
Data Read from External Filesystem	Data Read from IO Cache	Data Written to IO Cache	Estimated Size of Repeatedly Read Data	Cache Metadata Manager Peak Disk Usage		
260.9 GB	52.3 GB	260.9 GB	53.1 GB (16 %) - 53.1 GB (16 %)	4.3 KB		

Hot Data Auto Cached
Super Fast
Relieves Memory Pressure

```
spark.databricks.io.cache.enabled true  
spark.databricks.io.cache.maxDiskUsage 50g  
spark.databricks.io.cache.maxMetaCache 1g  
spark.databricks.io.cache.compression.enabled false
```

Spark cache vs. DBIO cache

Higher is better



Why So Fast?

HOW TO USE

AWS - i3s – On By Default
AZURE – Ls-series – On By Default

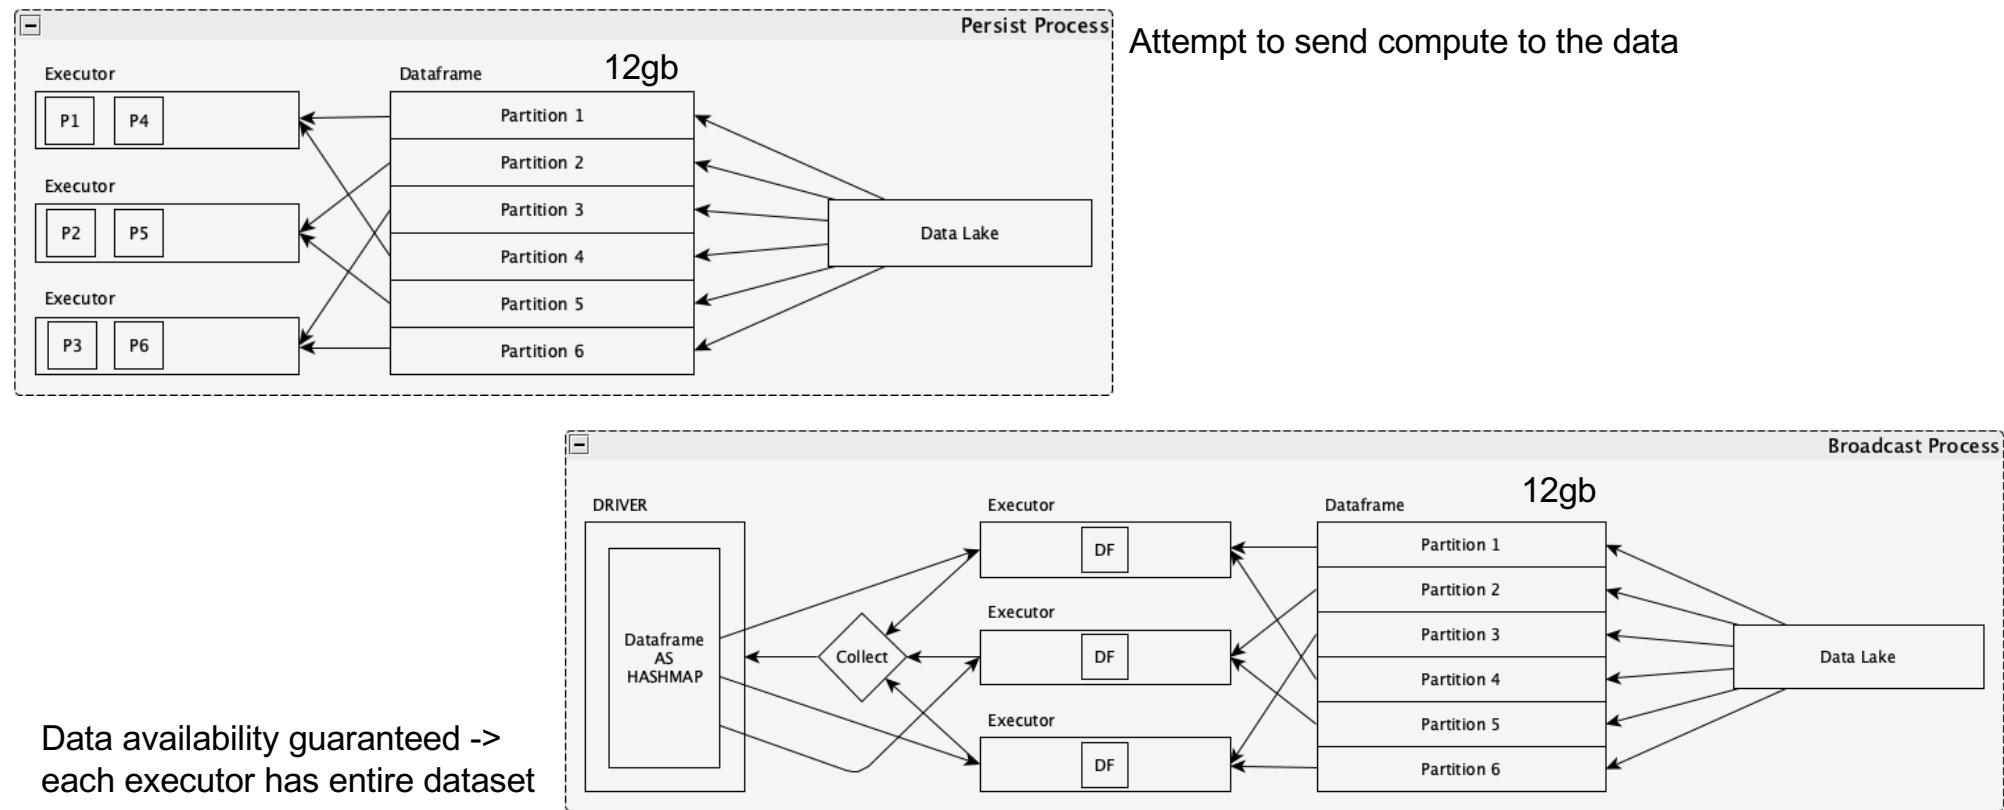
Join Optimization

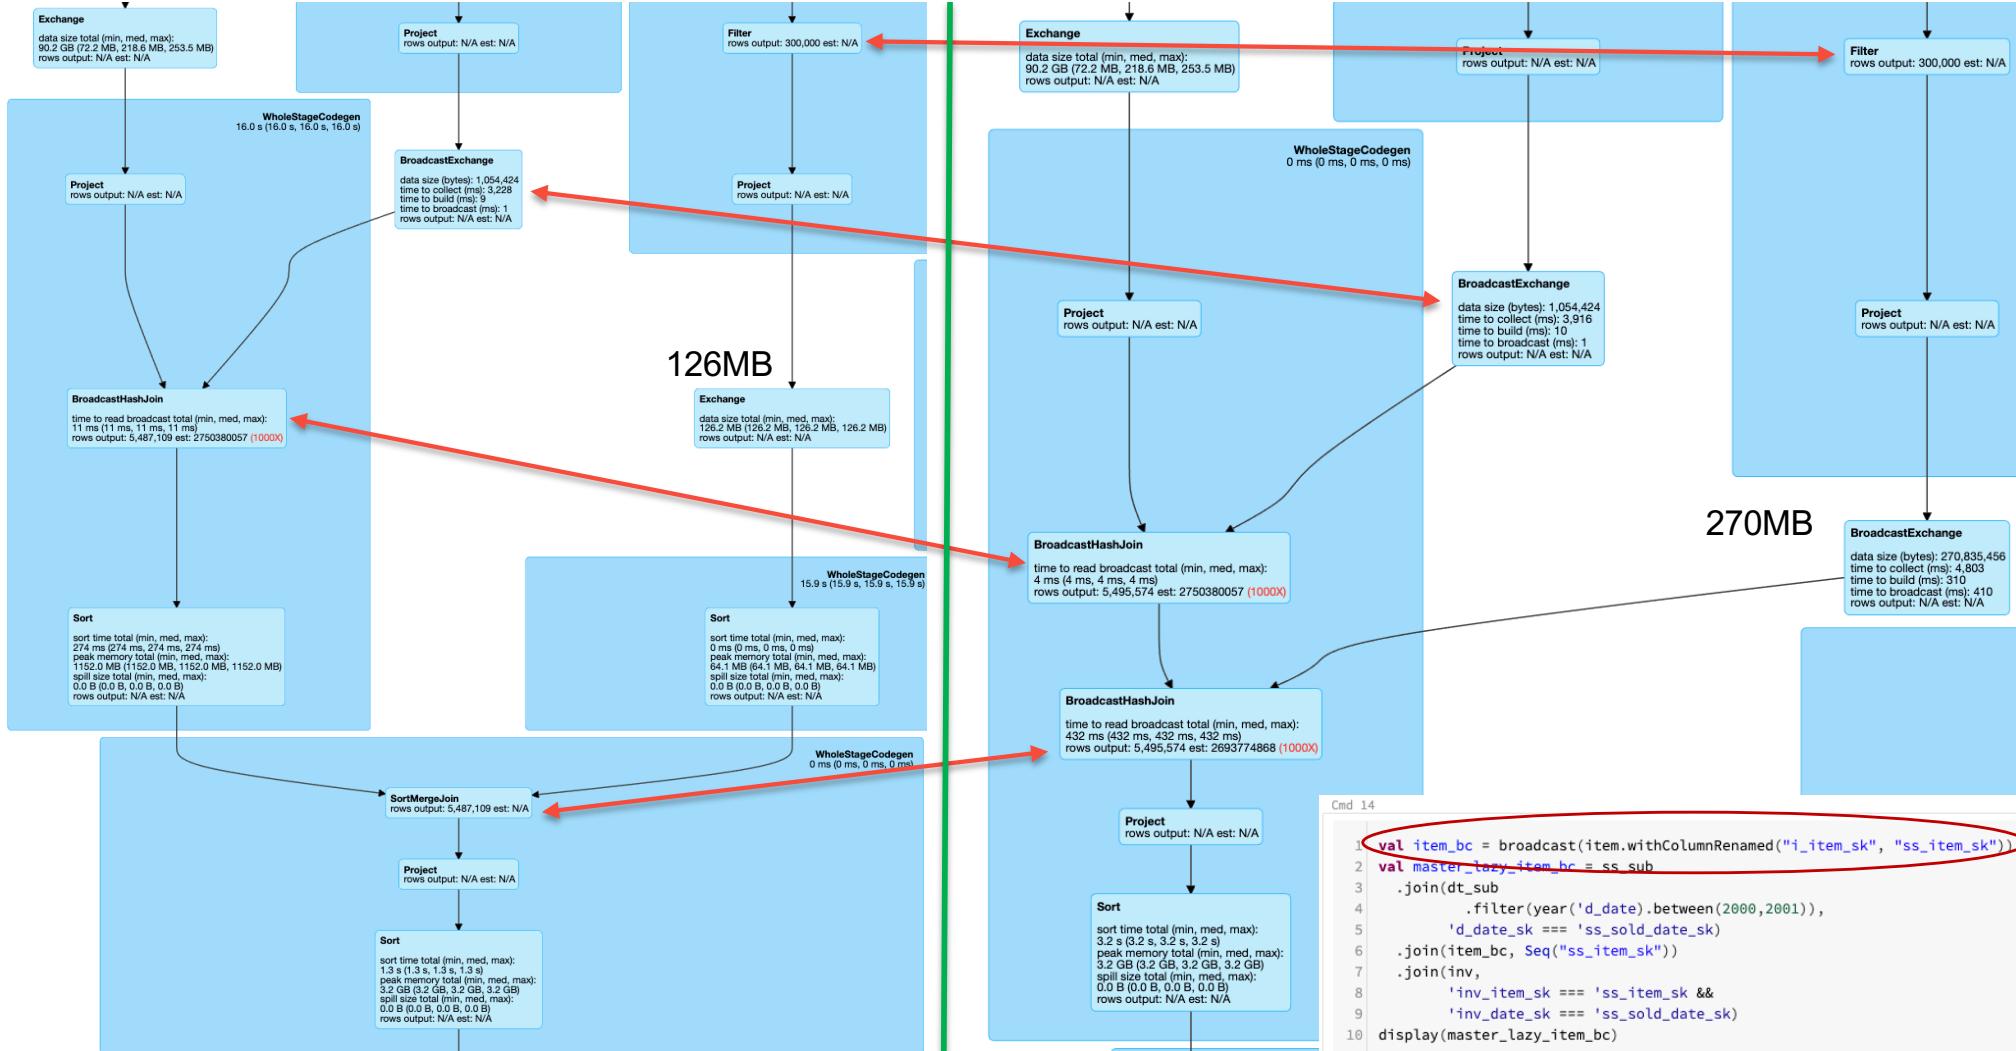
- SortMergeJoins (Standard)
- Broadcast Joins (Fastest)
- Skew Joins
- Range Joins
- BroadcastedNestedLoop Joins (BNLJ)

Join Optimization

- SortMerge Join – Both sides are large
- Broadcast Joins – One side is small
 - Automatic If:
(one side < *spark.sql.autoBroadcastJoinThreshold*) (default 10m)
 - Risks
 - Not Enough Driver Memory
 - DF > *spark.driver.maxResultSize*
 - DF > Single Executor Available Working Memory
 - Prod – Mitigate The Risks
 - Validation Functions

Persistence Vs. Broadcast





From 6h and barely started
 TO 8m → Lazy Loading
 TO 2.5m → Broadcast

Cmd 31 (+)

```

1 | val master_lazy = ss_sub
2 |   .join(dt_sub
3 |     .filter(year('d_date').between(2000,2001)),
4 |     'd_date_sk === 'ss_sold_date_sk')
5 |   .join(broadcast(item.withColumnRenamed("i_item_sk", "ss_item_sk")), Seq("ss_item_sk"))
6 |   .join(inv,
7 |     'inv_item_sk === 'ss_item_sk &&
8 |     'inv_date_sk === 'ss_sold_date_sk)

```

▼Completed Stages (3)

Stage Id	Pool Name	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
36	4260198380812558228	master_lazy.write.format("parquet").mode("overw... save at command-885896:1 +details")	2019/04/19 17:42:45	2.5 min	200/200	21.8 GB	55.6 GB		
35	4260198380812558228	master_lazy.write.format("parquet").mode("overw... save at command-885896:1 +details")	2019/04/19 17:42:02	43 s	452/452	14.7 GB			47.1 GB
34	4260198380812558228	master_lazy.write.format("parquet").mode("overw... save at command-885896:1 +details")	2019/04/19 17:42:02	21 s	261/261	3.8 GB			8.6 GB

SQL Example:

SELECT /*+ BROADCAST(customers) */ * FROM customers, orders WHERE o_custId = c_custId

Skew Join Optimization

Cmd 5

```
1 val skewedVals = invWSkew.withColumn("key", concat('inv_item_sk, 'inv_date_sk))
2   .groupBy("key", "inv_item_sk", "inv_date_sk")
3   .agg(count('key).alias("cnt"))
4   .orderBy('cnt.desc)
```

Cmd 6

display(skewedVals)

(1) Spark Jobs
Job 43 View (Stages: 2/2)
Stage 76: 229/229
Stage 77: 480/480

key	inv_item_sk	inv_date_sk	cnt
924052452275	92405	2452275	42872883
924042452275	92404	2452275	42857263
924072452275	92407	2452275	42857248
924062452275	92406	2452275	42854886
924032452275	92403	2452275	42853282
924012452275	92401	2452275	42852880
924022452275	92402	2452275	42851558
565032452341	56503	2452341	20
1196112452341	119611	2452341	20
2926412452341	292641	2452341	20
2283662452341	228366	2452341	20

- OSS Fix - Salting

- Add Column to each side with random int between 0 and spark.sql.shuffle.partitions – 1 to both sides
- Add join clause to include join on generated column above
- Drop temp columns from result

- Databricks Fix ([Skew Join](#))

```
val skewedKeys = List("id1", "id200", "id-99")
df.join(
  skewDF.hint("SKEW", "skewKey", skewedKeys),
  Seq(keyCol), "inner")
```

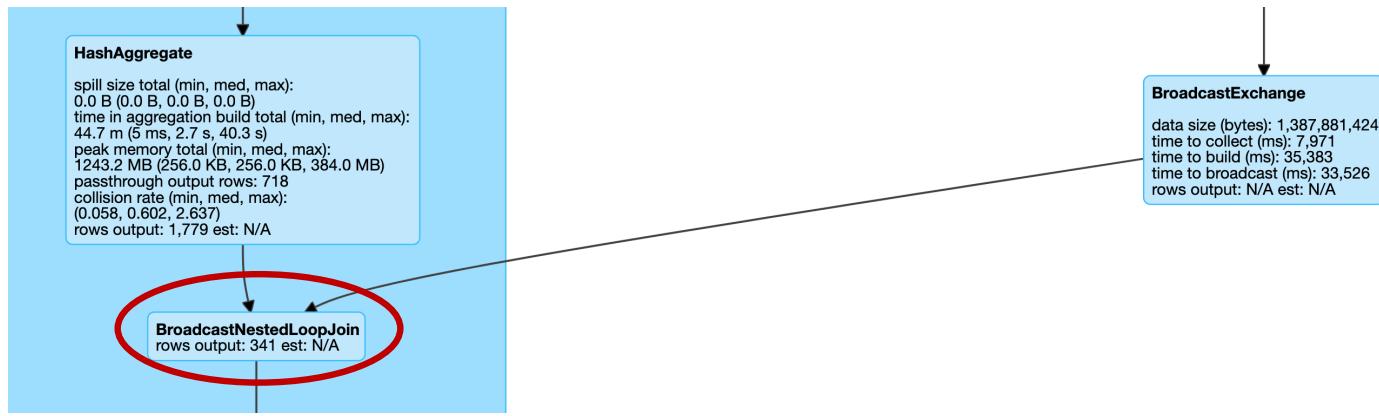
Skewed Aggregates

```
df.groupBy("city", "state").agg(<f(x)>).orderBy(col.desc)
```

```
val saltVal = random(0, spark.conf.get(org...shuffle.partitions) - 1)
```

```
df.withColumn("salt", lit(saltVal))  
    .groupBy("city", "state", "salt")  
    .agg(<f(x)>)  
    .drop("salt")  
    .orderBy(col.desc)
```

BroadcastNestedLoopJoin (BNLJ)



```

1  SELECT distinct sys_cd_val , PI_NUM FROM retail_sales
2  WHERE loadTS IN (SELECT MAX(loadTS) FROM retail_sales )
3  AND PI_NUM NOT IN (SELECT distinct PI_13 FROM product_master)
4  AND sys_cd_val <> 'PWG'

```

Stage Id	Pool Name	Description	Submitted	Duration	Tasks: Succeeded/total	Input	Output	Shuffle Read	Shuffle Write	Failure Reason
47	3925260204415598680	SELECT distinct SRC_SYS_CD , PIN_NUM FROM edl... collectResult at OutputAggregator.scala:136 +details	2019/04/22 18:56:06	7.1 min	557/563 (6)	7.8 GB				Job 17 cancelled part of cancelled job 3925260204415598680_50868365920

```

1  SELECT distinct sys_cd_val , PI_NUM
2  FROM retail_sales
3  WHERE loadTS IN (
4      SELECT MAX(loadTS) FROM retail_sales
5  )
6  AND NOT exists (
7      SELECT 1
8      FROM product_master sub
9      where sub.PI_13 <= a.PI_NUM
10 )
11 AND sys_cd_val <> 'PWG'

```

Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
019/04/22 19:04:34	0.1 s	1/1			7.1 KB	
019/04/22 19:04:06	28 s	200/200			674.0 MB	1062.5 KB
019/04/22 19:03:50	16 s	569/569	6.6 GB			148.8 MB
019/04/22 19:03:50	14 s	12/12				525.2 MB

Range Join Optimization

- Range Joins Types
 - Point In Interval Range Join
 - Predicate specifies value in one relation that is between two values from the other relation
 - Interval Overlap Range Join
 - Predicate specifies an overlap of intervals between two values from each relation

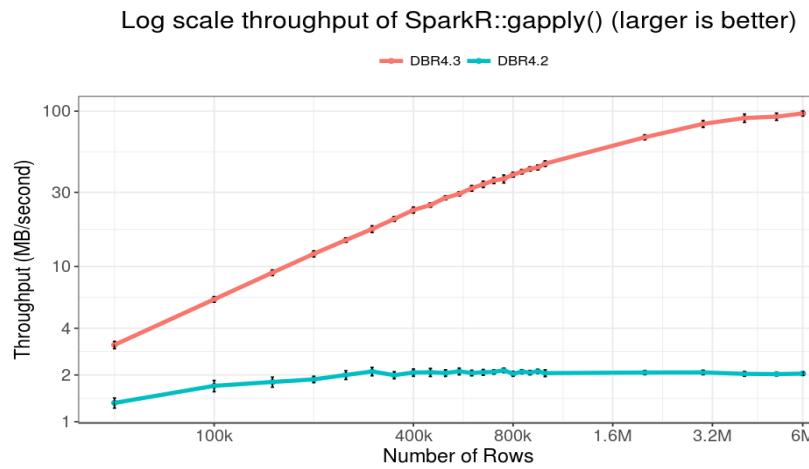
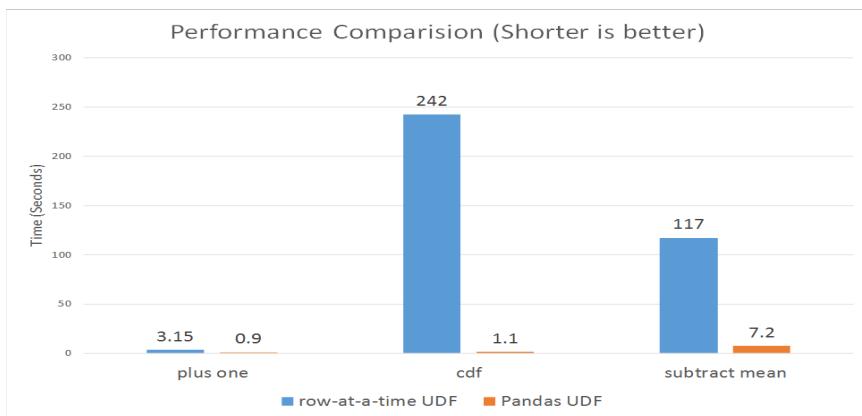
REFERENCE

Omit Expensive Ops

- Repartition
 - Use Coalesce or Shuffle Partition Count
- Count – Do you really need it?
- DistinctCount
 - use approxCountDistinct()
- If distincts are required, put them in the right place
 - Use dropDuplicates
 - dropDuplicates BEFORE the join
 - dropDuplicates BEFORE the groupBy

UDF Penalties

- Traditional UDFs cannot use [Tungsten](#)
 - Use [org.apache.spark.sql.functions](#)
 - Use [PandasUDFs](#)
 - Utilizes [Apache Arrow](#)
 - Use [SparkR UDFs](#)



Advanced Parallelism

- Spark's Three Levels of Parallelism
 - Driver Parallelism
 - Horizontal Parallelism
 - Executor Parallelism

```

1  val origHist_1y = spark.read.table(origHistTable)
2
3  spark.conf.set("spark.sql.shuffle.partitions", 500)
4
5  private val taskSupport = new ForkJoinTaskSupport(new ForkJoinPool(parallelism))
6  case class monthPart(year: Int, month: Int)
7
8  private val parts: ArrayBuffer[monthPart] = ArrayBuffer[monthPart]()
9  Range(2018,2019).toArray.foreach(yr => {
10    Range(1,13).toArray.foreach(mnth => parts.append(monthPart(yr, mnth)))
11  })
12
13  val parMonth = parts.toArray.par
14
15  parMonth.tasksupport = taskSupport
16  parMonth.foreach(part => {
17    try {
18      origHist_1y.filter('end_yr === part.year && 'end_mnth === part.month)
19        .write.format("delta")
20        .mode("append")
21        .partitionBy("end_cptr_dt")
22        .save(histPerfPath_1y)
23    } catch {
24      case e: Throwable => println(s"ERROR: Failed on ${part.year}-${part.month} with, $e")
25    }
26    println(s"Completed ${part.year}-${part.month}")
27  })

```

Summary

- Utilize Lazy Loading (Data Skipping)
- Maximize Your Hardware
- Right Size Spark Partitions
- Balance
- Optimized Joins
- Minimize Data Movement
- Minimize Repetition
- Only Use Vectorized UDFs

QUESTIONS